

The Exact Boundary of Period Inflation: A Machine-Checked Characterization for the Winnower Cellular-Automaton Background Detector

Joshua Herman¹

¹Independent Researcher
zitterbewegung@gmail.com

Abstract

A recurring task in artificial life is separating a life-like pattern from the periodic domain it lives in: a glider is a glider only relative to its background. *Winnower* is an open-source tool that performs this split on a cellular-automaton (CA) spacetime, fitting a drifting periodic *background* by majority vote and returning the *residual mask* of cells the background cannot explain. Like every fit-based detector it faces *period inflation*: a longer-period candidate has more parameters and can never fit worse, so a selector scoring by fit alone drifts to the largest model. We report a result that pins down exactly when this pressure exists. For two constant-velocity candidates, five natural conditions — velocity-matched divisibility, orbit-partition refinement, model-class nesting, universal residual improvement, and universal likelihood improvement — are all equivalent, and divisibility is *necessary*, not merely sufficient: whenever it fails, some spacetime makes the longer candidate strictly worse. The full equivalence, including the previously open necessity direction, is now machine-checked in Lean 4 under public continuous integration, with an exhaustive finite-window test as a cross-check. The characterization is the precise justification for Winnower’s complexity penalty, and it is reproducible and inspectable in the tool’s live web companion.

Introduction

Cellular automata routinely settle into broad periodic domains threaded with defects, interfaces, and particle-like activity. For artificial life those domains are not noise to be discarded but the *reference frame* that makes coherent structure legible: gliders, spaceships, and their collisions are defined against the background they traverse (Hanson and Crutchfield, 1992; Crutchfield and Hanson, 1997). Extracting that background at scale is a prerequisite for cataloguing where life-like behaviour appears — a concern squarely within the ALIFE 2026 theme of living and life-like complex adaptive systems.

Winnower is a fast global front end for this problem. Given a binary spacetime U and a grid of *candidates* (p, s) — “the pattern repeats every p steps, drifting sideways by s ” — it groups the cells each candidate forces to be equal (its *orbit classes*), takes a majority vote in each class to build the best-fitting background B^* , and records the residual mask

$M = U \oplus B^*$. Each candidate costs one linear pass, so whole-catalogue 1D/2D/3D surveys are practical, and every reported number re-runs in a browser-based reproduction.

A selector that scores candidates by fit alone carries a built-in bias. Call (p_2, s_2) a *velocity-matched multiple* of (p_1, s_1) when $p_2 = m p_1$ and $s_2 = m s_1 \bmod \mathbf{D}$: the same drift at a longer period. Such a multiple splits every orbit class into finer pieces, so its optimal residual count and its negative log-likelihood (NLL) can only improve. Unpenalized selection therefore inflates the period whether or not extra structure is present. Winnower counters this with a Bernoulli normalized-maximum-likelihood (NML) complexity penalty (Shtarkov, 1987; Grünwald, 2007), but a penalty is well-motivated only if we know *exactly* which comparisons carry the inflation pressure. This abstract reports that boundary and its machine-checked proof.

An Exact Boundary

Fix a torus with spatial dimensions $\mathbf{D}$ and temporal length T , and two candidates c_1, c_2 acting by spacetime translation. Write Π_i for the orbit partition of c_i , $\mathcal{B}_i$ for the model class of fields constant on Π_i -classes, $d_U^*(c_i)$ for the minimum Hamming distance from U to $\mathcal{B}_i$, and $\text{NLL}_U(c_i)$ for the orbit-class Bernoulli negative log-likelihood. The following are equivalent: (i) c_2 is a velocity-matched multiple of c_1 ; (ii) c_2 ’s translation is a power of c_1 ’s; (iii) Π_2 refines Π_1 ; (iv) $\mathcal{B}_1 \subseteq \mathcal{B}_2$; (v) $d_U^*(c_2) \leq d_U^*(c_1)$ for every U ; (vi) $\text{NLL}_U(c_2) \leq \text{NLL}_U(c_1)$ for every U .

The forward chain (i) $\Rightarrow$ $\dots$ $\Rightarrow$ (vi) is routine (cyclic-group structure; Jensen’s inequality for the concavity of binary entropy). The content is *necessity*: if divisibility fails, refinement fails, and a concrete adversarial spacetime — constant on Π_1 -classes but mixing two of them inside one Π_2 -class — makes c_2 strictly worse under both d^* and NLL. Closing the loop needs (iii) $\Rightarrow$ (ii), that refinement *forces* a power: on the torus orbits are cosets of the cyclic subgroup a translation generates, so refinement at the identity gives subgroup containment, and a translation is fixed by its value at one point. Divisibility is thus the exact frontier of the guaranteed-improvement regime — on one side longer pe-




Figure 1: Why the boundary matters in practice. On the 105 non-degenerate Life-like rules at horizon $T=100$, selecting by fit alone (raw residual count, or plain NLL) is pulled to the largest scanned period — the inflation the theorem predicts along every velocity-matched chain — while the complexity-penalized score (Bernoulli NML) concentrates at low period. Plain NLL selects the maximum period for *all* 105 rules.

riods always win, on the other the data can always punish them — and hence exactly the set of comparisons Winnower’s penalty must neutralize.

Machine-Checked, and Reproducible in the Tool

We treat a claim as verified only when Lean’s kernel accepts it. The full six-way equivalence, including the necessity direction (iii) $\Rightarrow$ (ii) that earlier drafts left open, is proved in Lean 4 against a pinned Mathlib toolchain and rebuilt on every push by the project’s public continuous-integration workflow, which additionally greps the built artifact for placeholder tactics. Independently, an exhaustive test enumerates all candidate pairs on a 6×4 space-time window and checks that refinement holds *iff* velocity-matched divisibility holds — a finite-window counterpart to the torus theorem and a guard against silent formalization drift.

The characterization is not a footnote to the tool; it is its justification, and it is inspectable where the tool lives. The Winnower web companion carries a dedicated late-breaking section stating the theorem in plain terms and linking the machine-checked Lean file and its CI runs; and the same in-browser reproduction that backs every other number reports, for any supported rule, seed, and horizon, the per-period fit (NLL) and penalized (NML) scores side by side — so a visitor can watch the fit column improve monotonically along velocity-matched chains, exactly as the theorem dictates, while the penalized column resists (fig. 1).

Status and Future Work

This is work in progress. The empirical picture — planted-ground-truth recovery at 1.00 joint accuracy over 642 runs,

zero false positives on shuffled and random controls, and the ablation above — is established in the companion full paper; the contribution here is the exact, machine-checked characterization explaining *why* the penalty has the shape it does and *exactly* where it is needed. Open threads include lifting the finite-window necessity direction (which currently carries one nondegeneracy hypothesis, $p_2 \leq T - 1$) into Lean alongside the torus form, and extending the analysis from pairwise comparisons to the full candidate lattice the scanner searches. The tool, the proofs, the tests, and the in-browser reproduction are public.

References

- Crutchfield, J. P. and Hanson, J. E. (1997). Computational mechanics of cellular automata: An example. *Physica D*, 103:169–189.
- Grünwald, P. (2007). *The Minimum Description Length Principle*. MIT Press.
- Hanson, J. E. and Crutchfield, J. P. (1992). The attractor-basin portrait of a cellular automaton. *Journal of Statistical Physics*, 66:1415–1462.
- Shtarkov, Y. (1987). Universal sequential coding of single messages. *Problems of Information Transmission*, 23:3–17.